

Beyond Prometheus: pushing the boundaries of scalable monitoring with VictoriaMetrics

Cloud Native Bergen | 28 October 2025

Diana Todea - DX Engineer

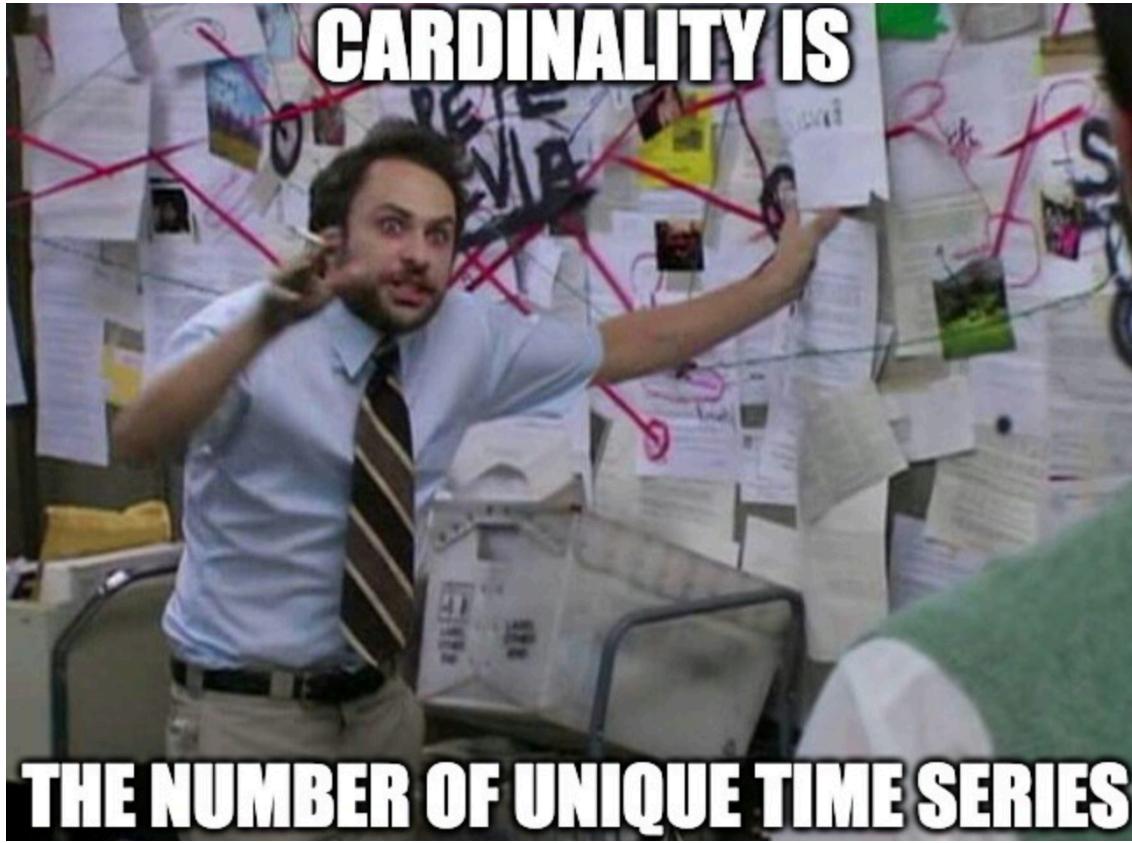
Bsky: @didiviking.bsky.social
X: @dianavtodea
Github: @didiViking/Conferences_Talks



Agenda

- ✓ Our common frenemy: cardinality
- ✓ To be or not to be Prometheus
- ✓ My next superhero
- ✓ Happy demo!
- ✓ Fatherly advice





HIGH CARDINALITY MEANS



HIGH NUMBER OF ACTIVE TIME SERIES



VictoriaMetrics

Simple, Reliable, Efficient Monitoring

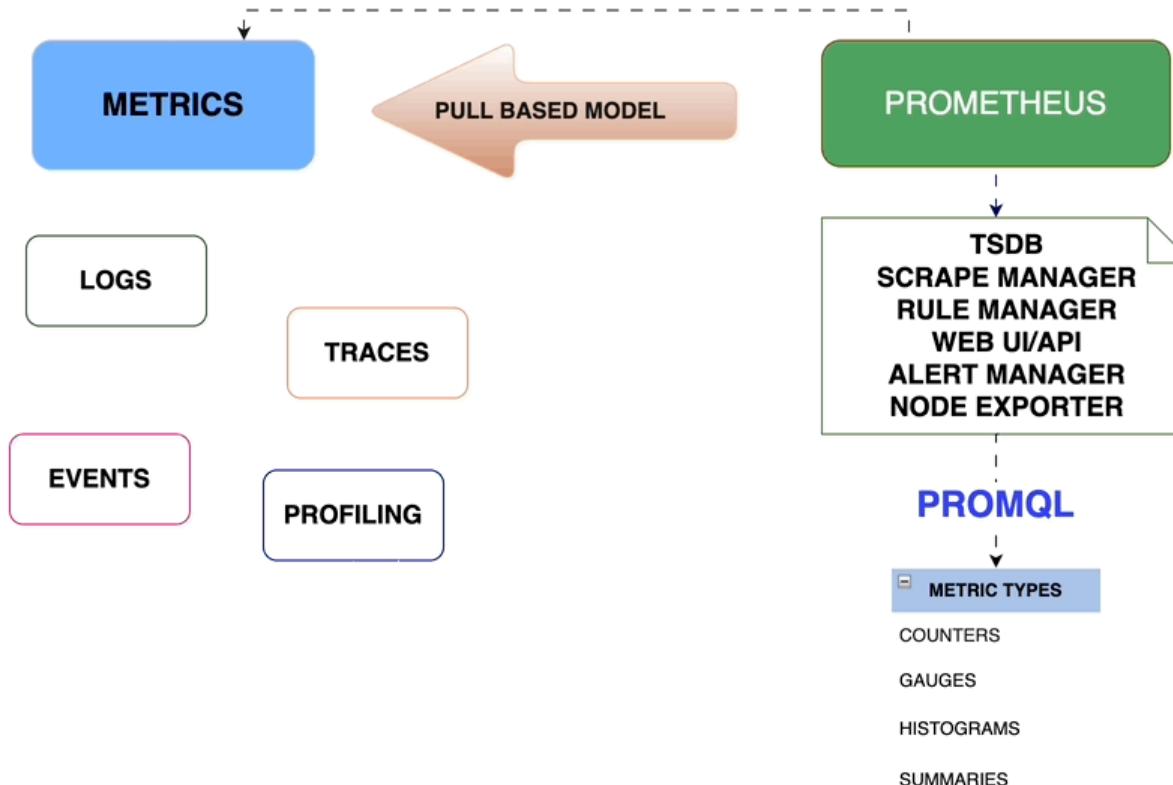
victoriametrics.com

High cardinality impact on o11y

- Increased resource consumption (storage, memory, CPU)
- Slower query performance
- System stability risks
- Increased cost
- Alert noise and signal dilution
- Increased complexity affecting actionable insights

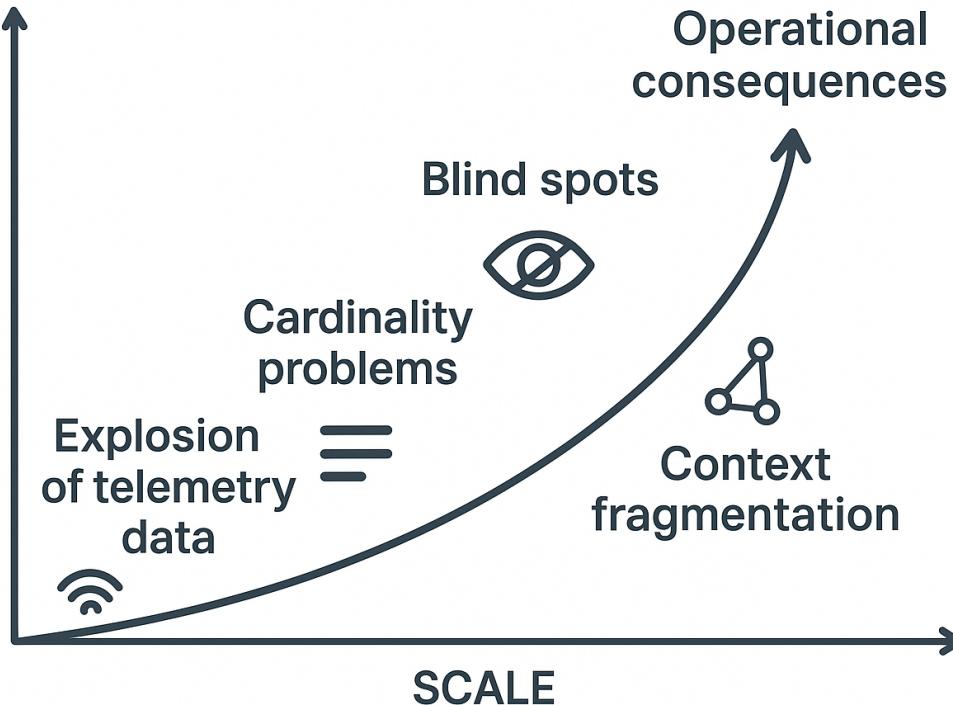


1 METRIC CAN HAVE MORE TIMESERIES



What happens when we scale?





High Cardinality = High Risk

- ☑ Pain points: out-of-memory errors, slow queries, alert storms

```
http_request_duration_seconds{user_id="12345", endpoint="/checkout", status="200", session_id="abcde123", pod="web-xyz"}
```



- ☑ user_id and session_id have huge cardinality

- ☑ every request generates a unique time series

- ☑ hundreds of thousands of active users

- ☑ hard to query, slow dashboards, OOM risk in Prometheus or alert manager



High cardinality root causes



labels with a large number of unique values (user_id, request_id, url, ip)



over-scraping/exporting



exporters that tag too much (unique IPs, device IDs)

Over-scraping

```
# Prometheus scrape config
scrape_configs:
  - job_name: 'node'
    scrape_interval: 5s
    static_configs:
      - targets: ['node1:9100', 'node2:9100']
```



You're over-collecting by 12x



Increased CPU/network usage



Your TSDB has to ingest/store redundant samples



Over-exporting

```
api_request_duration_seconds{endpoint="/search", status="200", user_id="1234", region="us-west", ...}
```

SOOO many labels

- ☑ Stop! Unique labels are killing your metrics
- ☑ Unused metrics = wasted resources
- ☑ Cardinality up, performance down



Over-tagging

```
node_disk_reads_total{device="nvme0n1"}  
node_disk_reads_total{device="nvme1n1"}  
node_disk_reads_total{device="nvme2n1"}
```



If all node exporter collectors are enabled each metric gets multiplied

Prometheus' limitations

- ☑ **Cardinality:** optimize with good labels, label drop, set limits on scrape jobs.
- ☑ **Long-term storage:** improve the architecture through **sharding, federation, and highly available replicas.**

Optimize Prometheus



Manage your **scrape jitters**, they are the **most common** cause of oversize TSDB blocks.

```
$ ./oy-scrape-jitter \
--prometheus.url=http://localhost:9090 \
--log.results-only
level=info aligned_targets=2 unaligned_targets=22 max_ms=30

prometheus:
prometheusSpec:
additionalArgs:
- name: "scrape.timestamp-tolerance"
value: "30ms"

$ ./oy-scrape-jitter \
--prometheus.url=http://localhost:9090 \
--log.results-only
level=info aligned_targets=23 unaligned_targets=1 max_ms=34
```



Optimize Prometheus



Use **pprof**, a GO native tool to check in detail Prometheus performance.

```
(pprof) top
Showing nodes accounting for 82.47MB, 67.76% of 121.71MB total
Dropped 137 nodes (cum <= 0.61MB)
Showing top 10 nodes out of 123
flat flat% sum% cum cum%
17.01MB 13.98% 13.98% 17.01MB 13.98% github.com/prometheus/
prometheus/model/labels.(*Builder).Labels
11.44MB 9.40% 23.37% 11.44MB 9.40% github.com/prometheus/
prometheus/scrape.(*scrapeCache).addRef
10.89MB 8.94% 32.32% 10.89MB 8.94% github.com/prometheus/
prometheus/scrape.newScrapePool.func1
9MB 7.40% 39.71% prometheus/tsdb/chunkenc.NewXORChunk
9MB 7.40% github.com/prometheus/
7.74MB 6.36% 46.08% Encode
7.74MB 6.36% github.com/golang/snappy.
7.71MB 6.33% 52.41% 7.71MB 6.33% github.com/prometheus/
prometheus/tsdb/index.(*MemPostings).addFor
6MB 4.93% 57.34% prometheus/tsdb.newMemSeries
9MB 7.40% github.com/prometheus/
5.19MB 4.26% 61.60% 5.19MB 4.26% github.com/prometheus/
prometheus/scrape.(*scrapeCache).trackStaleness
4MB 3.29% 64.89% (inline)
4MB 3.29% bufio.NewReaderSize
3.50MB 2.88% 67.76% 7.50MB 6.16% github.com/prometheus/
prometheus/tsdb.(*memSeries).mmapCurrentHeadChunk
```



Optimize Prometheus



Use **promtool** for pprof data: run **promtool debug all** before restarting Prometheus.

```
promtool debug pprof http://localhost:9090
collecting: http://localhost:9090/debug/pprof/threadcreate
collecting: http://localhost:9090/debug/pprof/profile?seconds=30
collecting: http://localhost:9090/debug/pprof/block
collecting: http://localhost:9090/debug/pprof/goroutine
collecting: http://localhost:9090/debug/pprof/heap
collecting: http://localhost:9090/debug/pprof/mutex
collecting: http://localhost:9090/debug/pprof/trace?seconds=30
Compiling debug information complete, all files written in "debug.tar.gz".
```

Optimize Prometheus

- ✓ Use recording rules: enabling the pre-computation of the expensive queries.

```
groups:  
- name: k8s.rules  
rules:  
- expr: |-  
    sum by (cluster, namespace, pod, container) (  
        rate(container_cpu_usage_seconds_total{job="kubelet",  
            metrics_path="/metrics/cadvisor", image!=""}[5m])  
    ) * on (cluster, namespace, pod) group_left(node) topk by  
        (cluster, namespace, pod) (1, max by(cluster, namespace, pod, node) (kube_pod_  
            info{node!=""}))  
    )  
    record: node_namespace_pod_container:container_cpu_usage_  
        seconds_total:sum_irate
```

Optimize Prometheus



Query logging is configured via the Prometheus configuration file under the global key with the `query_log_file` setting.

Each entry in this query log helps us find useful statistics about the query.

```
{  
  "params": {  
    "end": "2025-10-14T02:49:53.242Z",  
    "query": "max_over_time(reloader_last_reload_successful{namespace=~\".+\\"}[5m]) == 0",  
    "start": "2025-10-14T02:49:53.242Z",  
    "step": 0  
  },  
  "ruleGroup": {  
    "file": "/etc/prometheus/rules/rules_file.yaml",  
    "name": "config-reloaders"  
  },  
  "spanID": "0000000000000000",  
  "stats": {  
    "timings": {  
      "evalTotalTime": 0.000173986,  
      "resultSortTime": 0,  
      "queryPreparationTime": 0.000103703,  
      "innerEvalTime": 0.000063172,  
      "execQueueTime": 0.000017709,  
      "execTotalTime": 0.000197365  
    },  
    "samples": {  
      "totalQueryableSamples": 0,  
      "peakSamples": 2  
    }  
  },  
  "ts": "2025-10-14T02:49:53.243Z"  
}
```



Optimize Prometheus



Query limits: tune how many queries can be executed simultaneously and how broad those queries are allowed to be with command-line flags passed to the Prometheus at startup.

`--query.max-concurrency`

`--query.max-samples`



Optimize Prometheus

- ✓ **Tune garbage collection:** the Go garbage collector (GC) setting a percentage of how much the live heap size (memory) can increase from the previous garbage collection before the GC kicks in to mark and reclaim memory.

Optimize Prometheus

- ✓ Use the **GOMEMLIMIT** setting: it does not supersede GOGC and can be used in conjunction with it. It functions by providing a soft limit on the maximum amount of memory a Go program can safely use.



SHE'S LOOKING AT YOU, KID



VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

VictoriaMetrics Solutions

Open
Source

Metrics Logs
Traces

Enterprise
On Premise | Cloud |
Hybrid



VictoriaMetrics Key Features

Compatibility with
Prometheus and
Grafana

Compatibility with
Prometheus ecosystem:
Exporters, service
discovery, alerting rules,
query frameworks

Native Push/Pull
support without
limitations

Rich features set

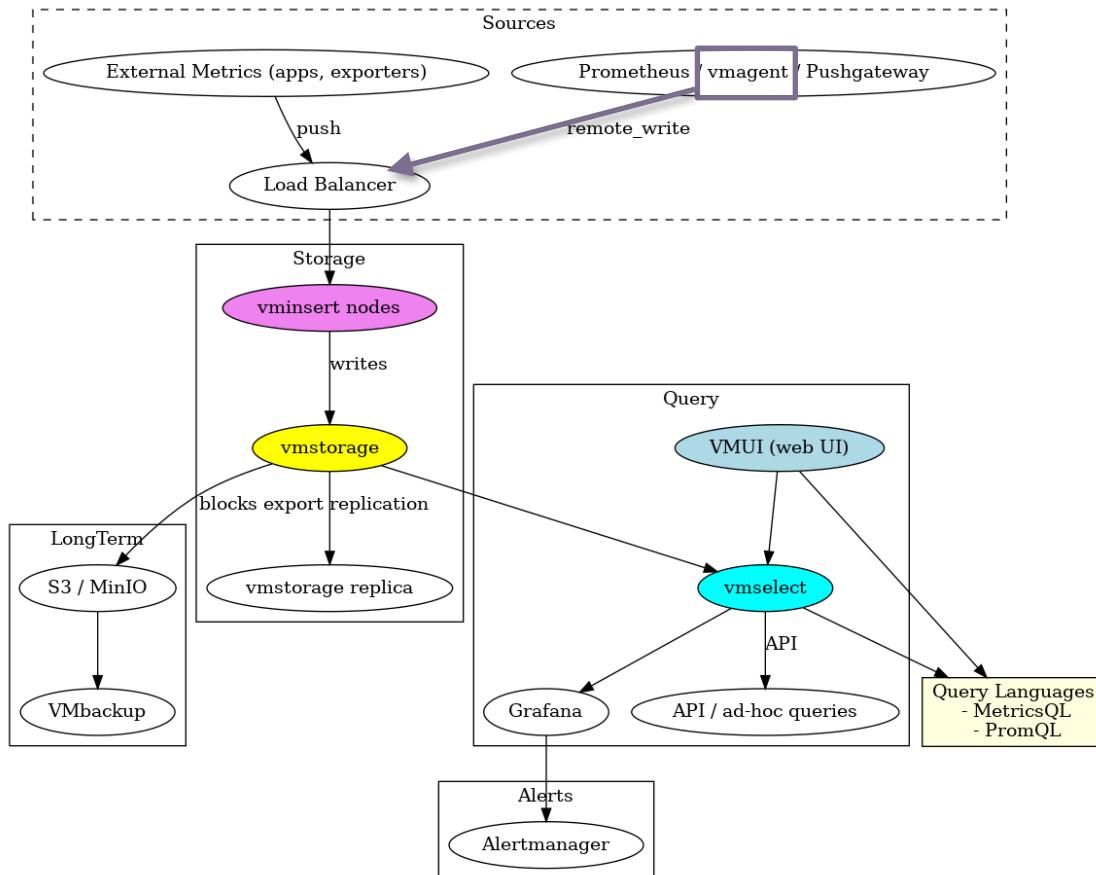
Clustering support
with multi-tenancy

Lightweight metrics
collector, best you
can find

Backup/restore
utilities

Auth
proxy



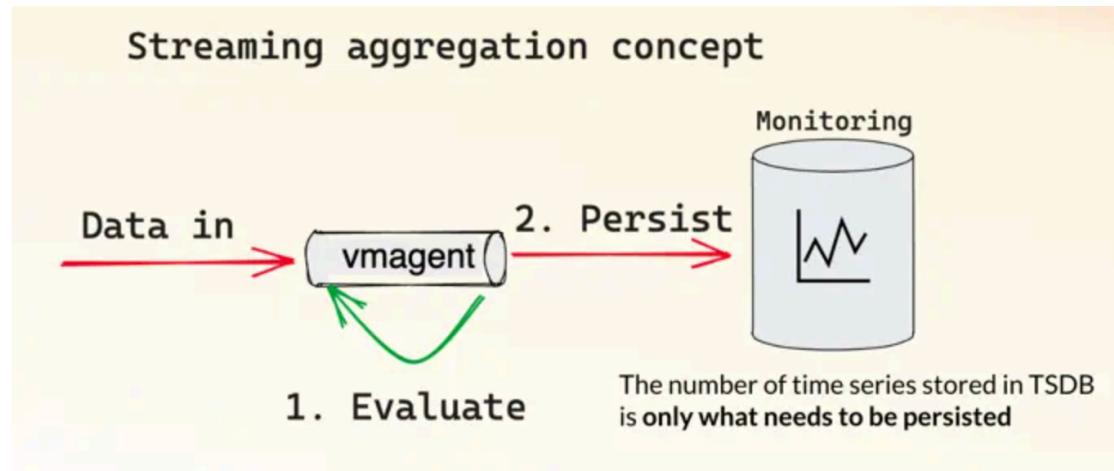


Main gains over Prometheus

Improved
benchmarks

	Prometheus	VictoriaMetrics
CPU avg used	0.79/3 cores	0.76/3 cores
Disk usage	83.5 GiB	✓ 33 GiB
Memory max used	8.12/12 GiB	✓ 4.5/12 GiB
Read latency 50th	70.5ms	✓ 4.3ms
Read latency 99th %ile	7s	✓ 3.6s

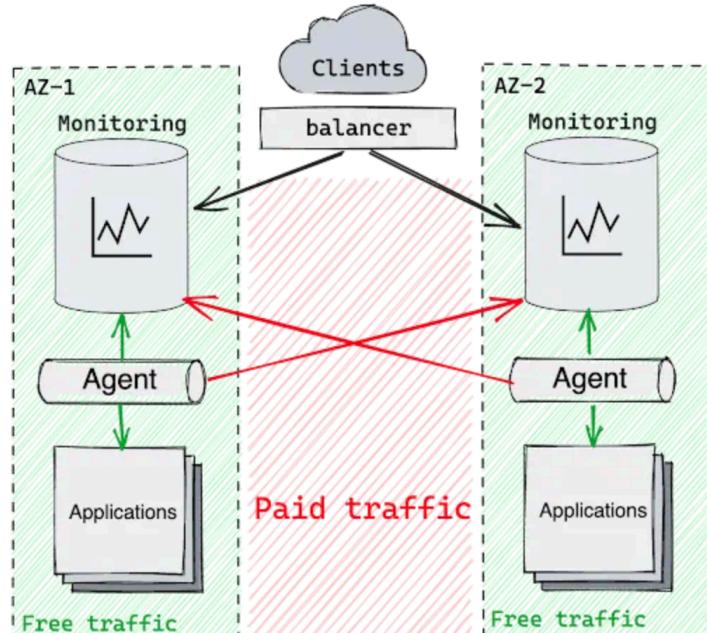
Main gains over Prometheus



Main gains over Prometheus

Save network
cost with
vmagent

Settings	Trade-off
remoteWrite.vmProtoCompressLevel	Increased compression level, traded for higher CPU usage
remoteWrite.maxBlockSize, remoteWrite.maxRowsPerBlock, remoteWrite.flushInterval	Increased batch size leading to better compression ratio, traded for latency
remoteWrite.significantFigures, remoteWrite.roundDigits	Reduced precision/entropy, better compression ratio



Cardinality explorer helps identify

 Metric names with the highest number of series

.....

 Labels with the highest number of series

.....

 Values with the highest number of series for the selected label

.....

 label=name pairs with the highest number of series

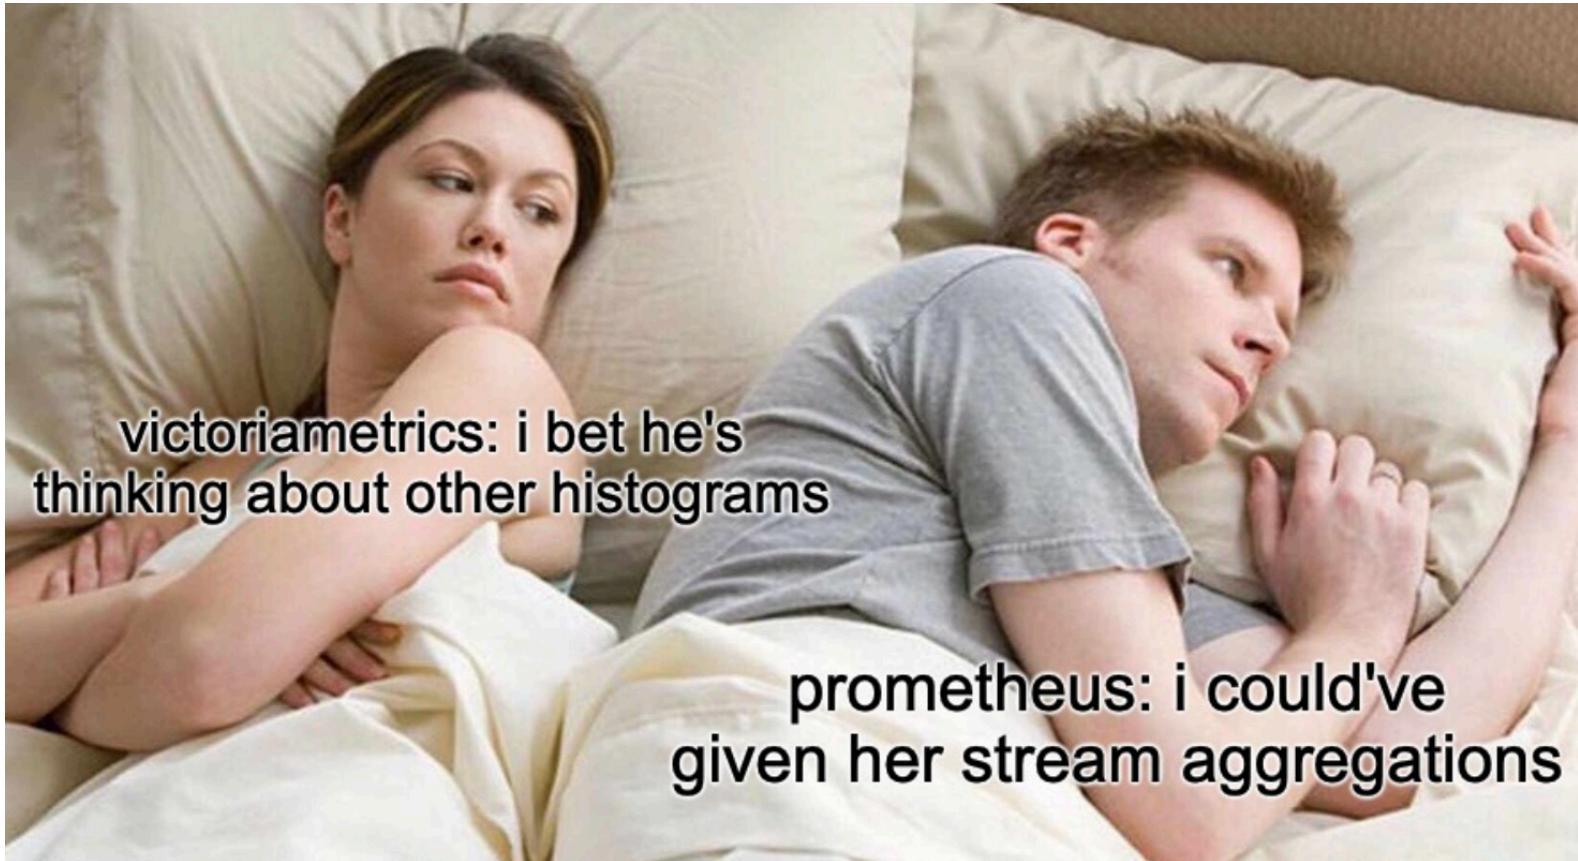
.....

 Labels with the highest number of unique values

.....



Nobody:



victoriometrics: i bet he's
thinking about other histograms

prometheus: i could've
given her stream aggregations

VictoriaMetrics Query Language



MetricsQL

- ☑ PromQL compatible, no need to change dashboards or alerting rules
- ☑ 100+ additional functions to improve querying experience
- ☑ Templating support in queries

Top 5 features

- ☑ Graphite compatible filters and Graphite support
- ☑ group_left(*) / group_right(*) with prefix support
- ☑ @ modifier for arbitrary timestamps and subexpressions
- ☑ default / if / ifnot binary operators
- ☑ WITH templates + string concatenation

Spot high cardinality on VictoriaMetrics clusters



High memory usage on vmstorage – indexes for many unique series sit in memory

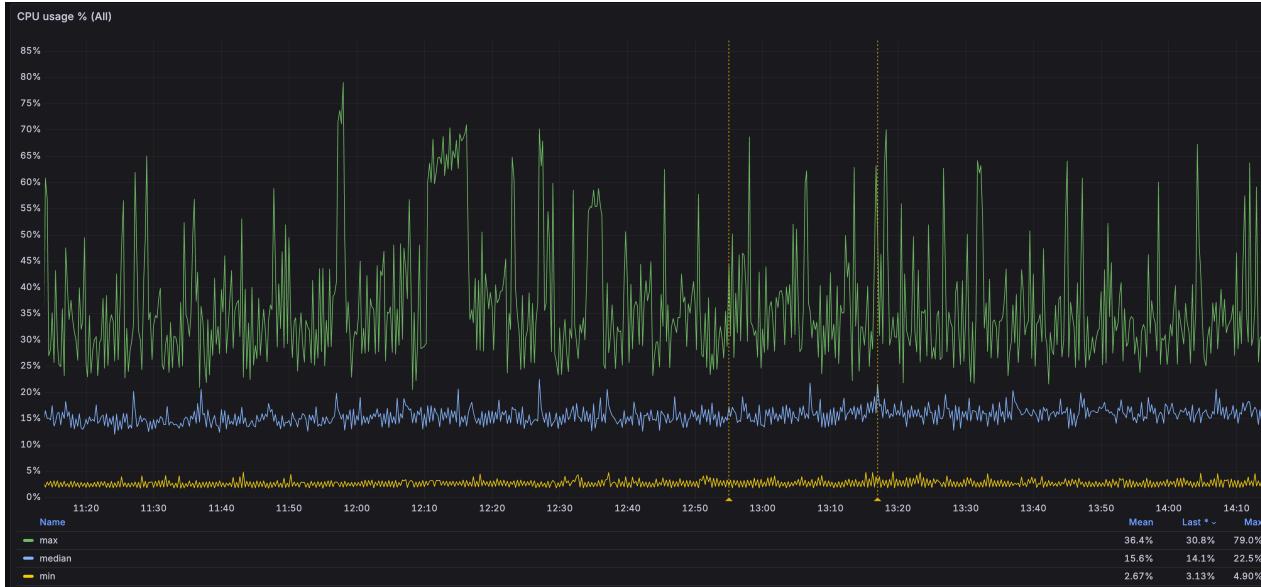


VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

vmstorage



CPU usage spike on vmstorage, each unique series requires a separate TSID mapping and index entry



vmselect - query processing



Memory spikes on vmselect, when queries need to touch millions of series

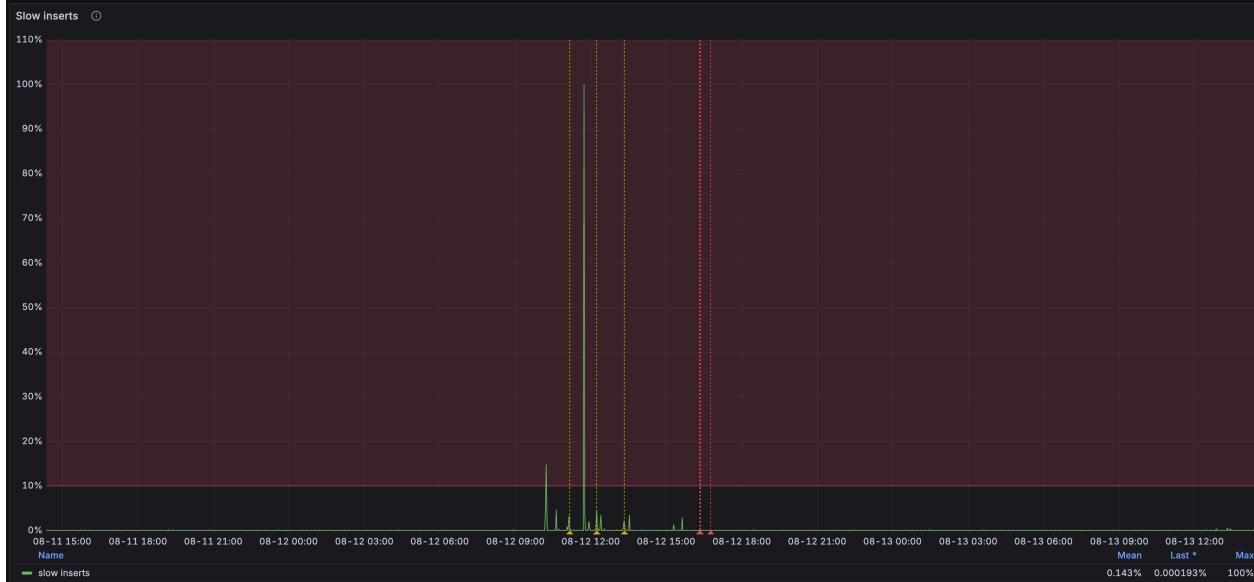


VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

vminsert connects to multiple vmstorage nodes



Increased ingestion latency, vminsert takes longer to batch and deduplicate incoming series before writing

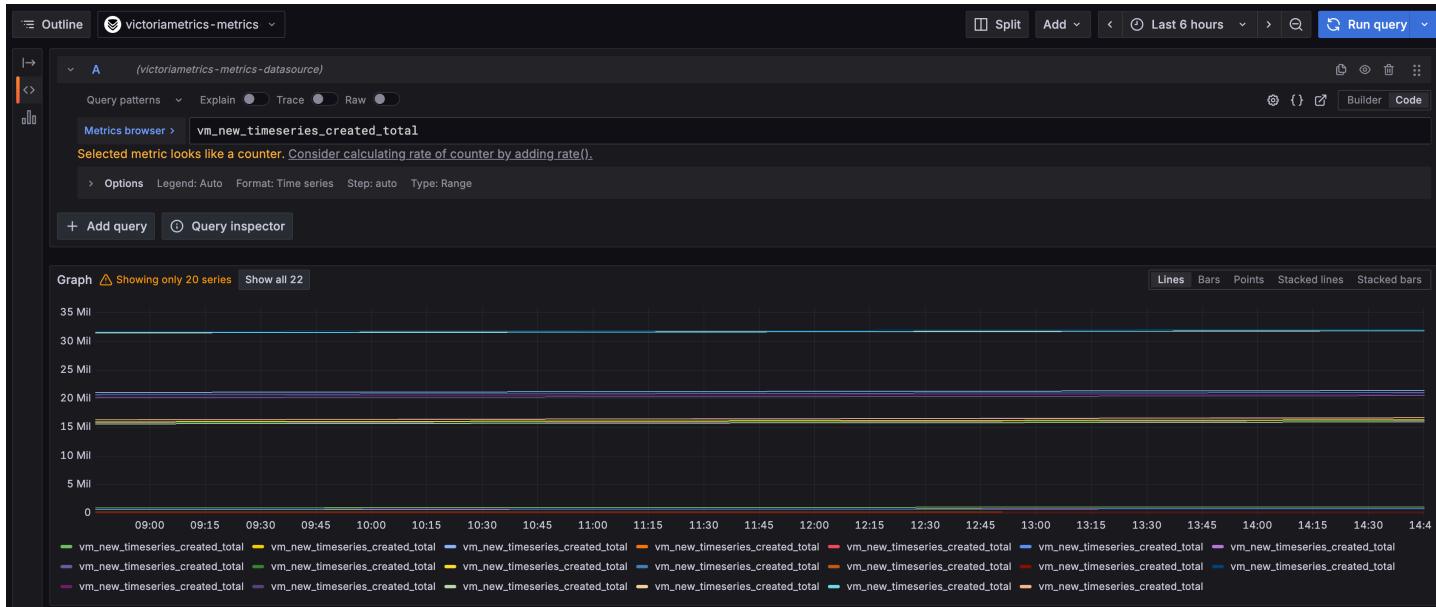


VictoriaMetrics

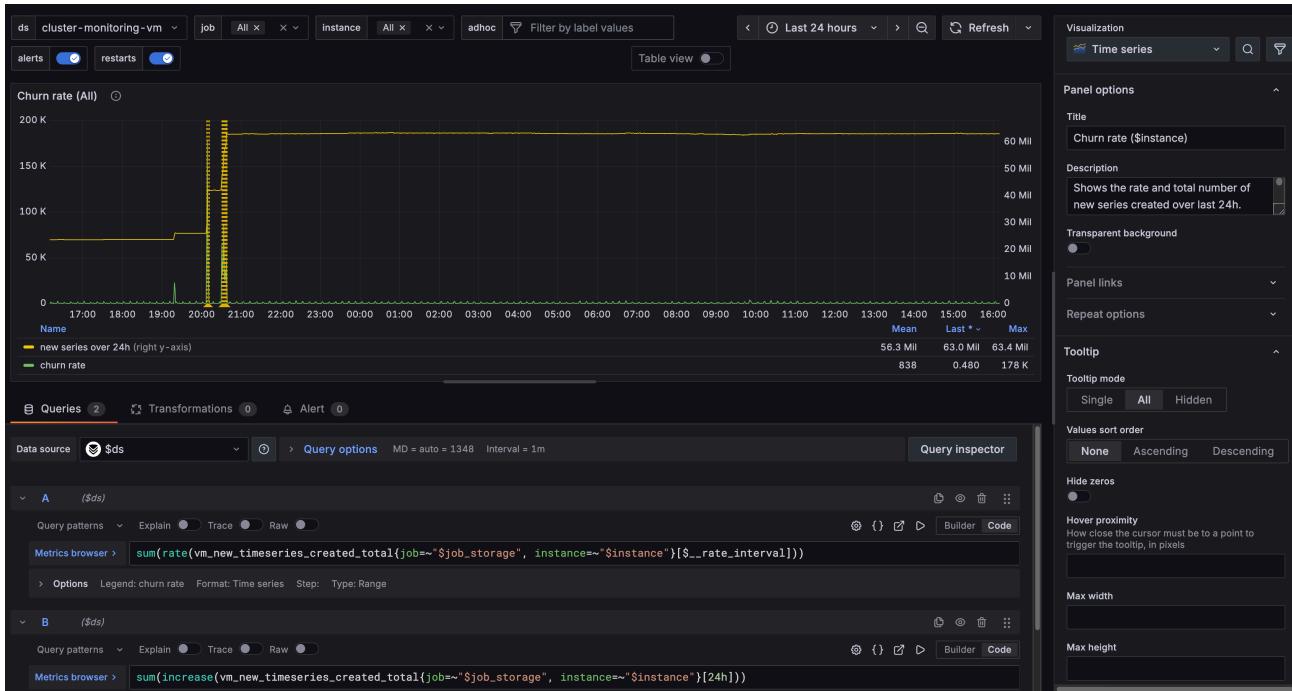
Simple, Reliable, Efficient Monitoring

victoriametrics.com

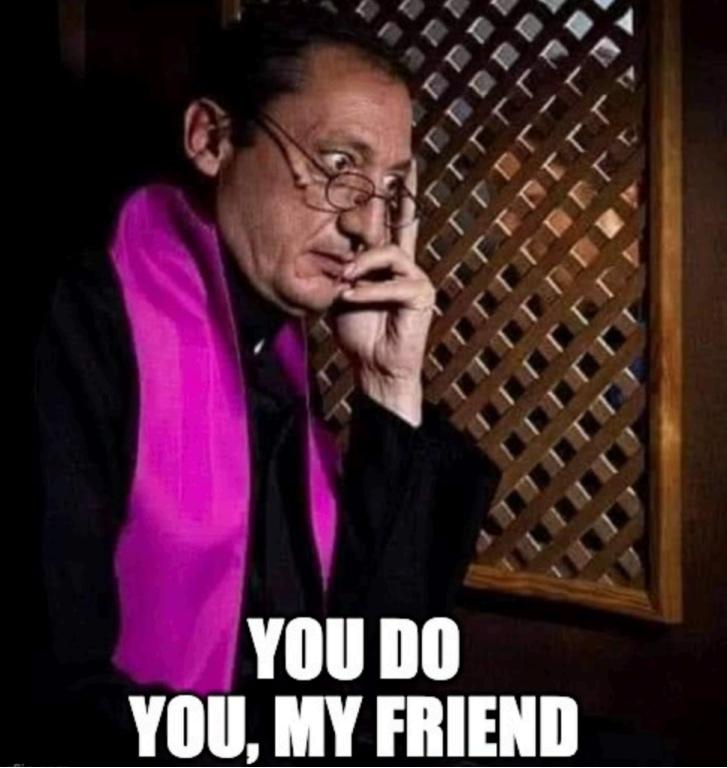
Useful queries



How to use it in a Grafana dashboard



WHAT SHOULD I DO, FATHER?



YOU DO YOU, MY FRIEND



VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

Take-aways

- ✓ Reduce histogram buckets, aggregate data with functions like sum() or avg()
- ✓ Look for metrics insights in vmui's Cardinality Explorer
- ✓ If labels can't be removed, pre-aggregate data before ingestion with stream aggregation
- ✓ Reduce query load by aggregating with MetricsQL instead of pulling raw series
- ✓ VictoriaMetrics uses optimized inverted indexes, efficient storage, aggressive compression

Resources

- ✓ <https://play.victoriametrics.com/> & <https://play-grafana.victoriametrics.com/>
- ✓ <https://github.com/VictoriaMetrics/prometheus-benchmark>
- ✓ <https://victoriametrics.com/blog/cardinality-explorer/>
- ✓ <https://medium.com/@romanhavronenko/victoriametrics-promql-compliance-d4318203f51e>
- ✓ <https://valyala.medium.com/prometheus-vs-victoriametrics-benchmark-on-node-exporter-metrics-4ca29c75590f>
- ✓ <https://github.com/VictoriaMetrics-Community/opentelemetry-demo>

Explore VictoriaMetrics Solutions



VictoriaMetrics

Open Source Available on Github



VictoriaTraces

Open Source Available on Github



VictoriaLogs

Open Source Available on Github



Tussen takk

Spørsmål?

Bsky: @didiviking.bsky.social

X: @dianavtodea

Github: @didiViking/Conferences_Talks

LinkedIn: @diana-todea-b2a79968



VictoriaMetrics Community