

Beyond Prometheus: pushing the boundaries of scalable monitoring with VictoriaMetrics

Criteo, Paris | 15 October 2025

Diana Todea - DX Engineer

Bsky: @didiviking.bsky.social
X: @dianavtodea
Github: @didiViking/Conferences_Talks



Agenda

Observability

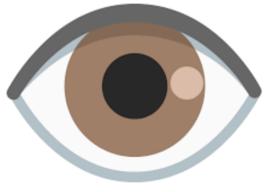
Prometheus

Cardinality

Demo

Practical
tips





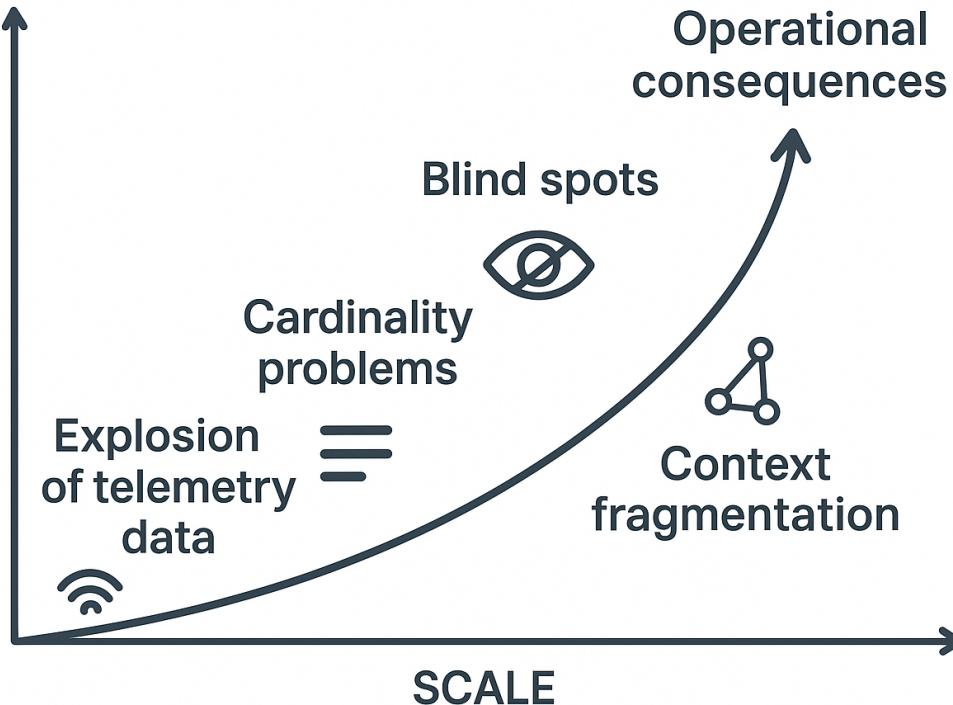
+



=

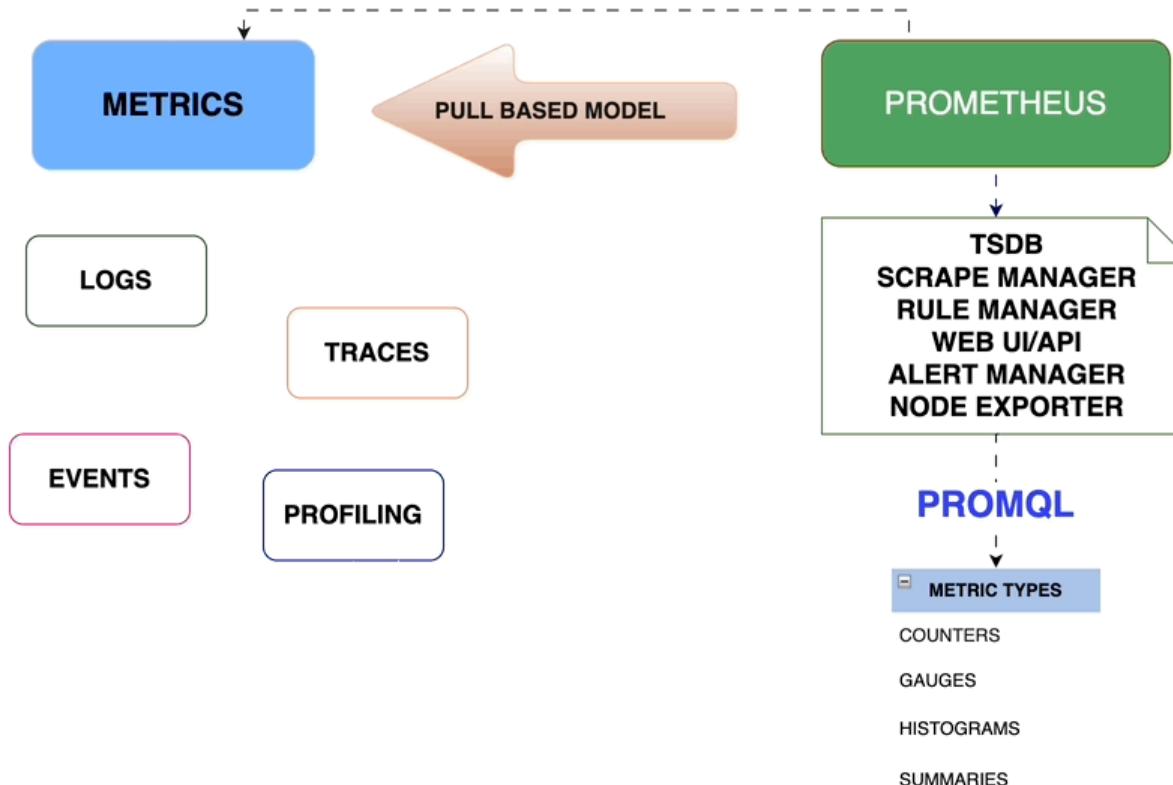
Actionable
Insights

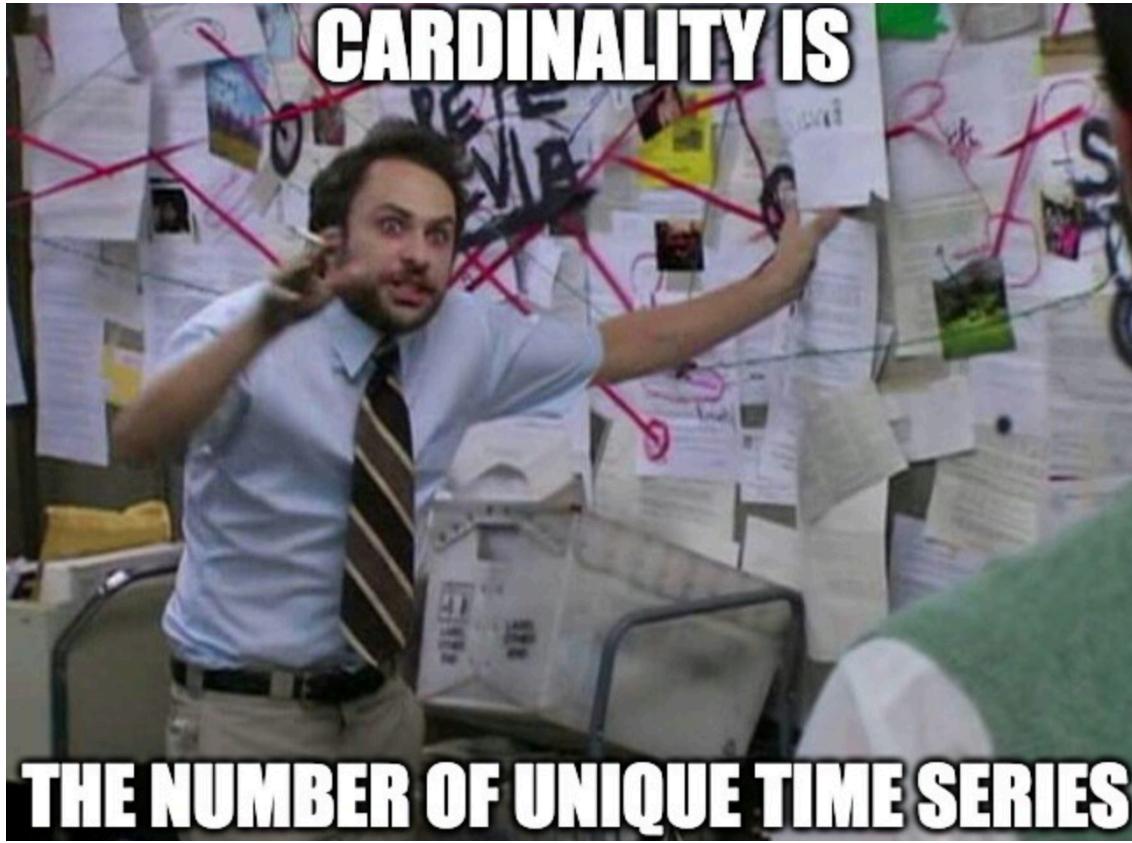
What happens when we scale?



Collect the right observability signals

1 METRIC CAN HAVE MORE TIMESERIES





HIGH CARDINALITY MEANS



HIGH NUMBER OF ACTIVE TIME SERIES



VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

High cardinality impact on o11y

- Increased resource consumption (storage, memory, CPU)
- Slower query performance
- System stability risks
- Increased cost
- Alert noise and signal dilution
- Increased complexity affecting actionable insights



High Cardinality = High Risk



Pain points: out-of-memory errors, slow queries, alert storms

```
http_request_duration_seconds{user_id="12345", endpoint="/checkout", status="200", session_id="abcde123", pod="web-xyz"}
```



user_id and session_id have huge cardinality



every request generates a unique time series



hundreds of thousands of active users



hard to query, slow dashboards, OOM risk in Prometheus or alert manager



High cardinality root causes



labels with a large number of unique values (user_id, request_id, url, ip)



over-scraping/exporting



exporters that tag too much (unique IPs, device IDs)

Over-scraping

```
# Prometheus scrape config
scrape_configs:
  - job_name: 'node'
    scrape_interval: 5s
    static_configs:
      - targets: ['node1:9100', 'node2:9100']
```



You're over-collecting by 12x



Increased CPU/network usage



Your TSDB has to ingest/store redundant samples



Over-exporting

```
api_request_duration_seconds{endpoint="/search", status="200", user_id="1234", region="us-west", ...}
```

SOOO many labels

- ☑ Stop! Unique labels are killing your metrics
- ☑ Unused metrics = wasted resources
- ☑ Cardinality up, performance down

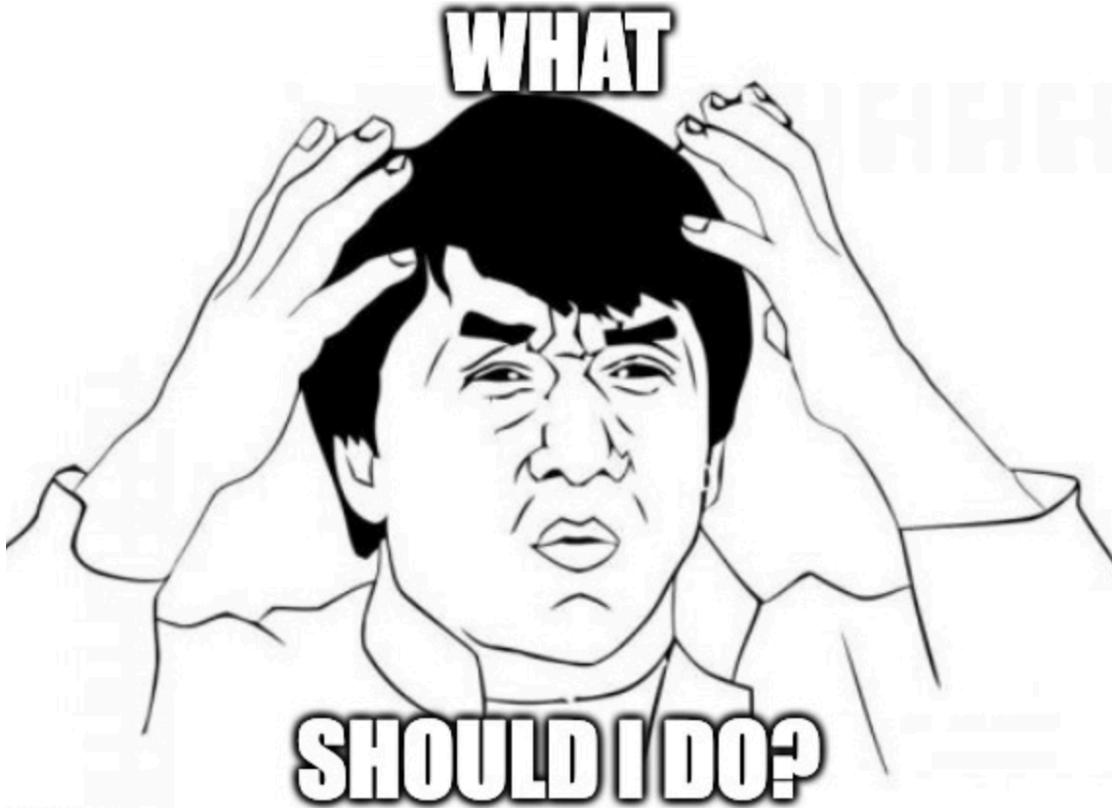


Over-tagging

```
node_disk_reads_total{device="nvme0n1"}  
node_disk_reads_total{device="nvme1n1"}  
node_disk_reads_total{device="nvme2n1"}
```



If all node exporter collectors are enabled each metric gets multiplied



Scrape strategies

- ✓ Don't scrape faster than metrics change

```
scrape_interval: 60s
```

- ✓ Drop unneeded labels

```
relabel_configs:  
  - source_labels: [__name__, user_id]  
    regex: api_request_duration_seconds;.*  
    action: labeldrop
```

Use pre-aggregated views

✓ Aggregate early and export less (stream aggregation)

✓ Summarize with histograms, not IDs



```
api_request_duration_seconds_bucket{le="0.1", endpoint="/search"}
```

Instance grouping & job deduplication

```
- job_name: app-v1  
- job_name: app-v2
```



```
- job_name: app  
  relabel_configs:  
    - source_labels: [__meta_kubernetes_label_app_version]  
      target_label: version
```

Cut the noise at the source

```
metric_relabel_configs:  
  - source_labels: [__name__]  
    regex: "go_gc_.*"  
    action: drop
```

(*if not used)

VictoriaMetrics Solutions

Open
Source

Metrics Logs
Traces

Enterprise
On Premise | Cloud |
Hybrid



VictoriaMetrics Key Features

Compatibility with
Prometheus and
Grafana

Compatibility with
Prometheus ecosystem:
Exporters, service
discovery, alerting rules,
query frameworks

Native Push/Pull
support without
limitations

Rich features set

Clustering support
with multi-tenancy

Lightweight metrics
collector, best you
can find

Backup/restore
utilities

Auth
proxy



VictoriaMetrics Query Language

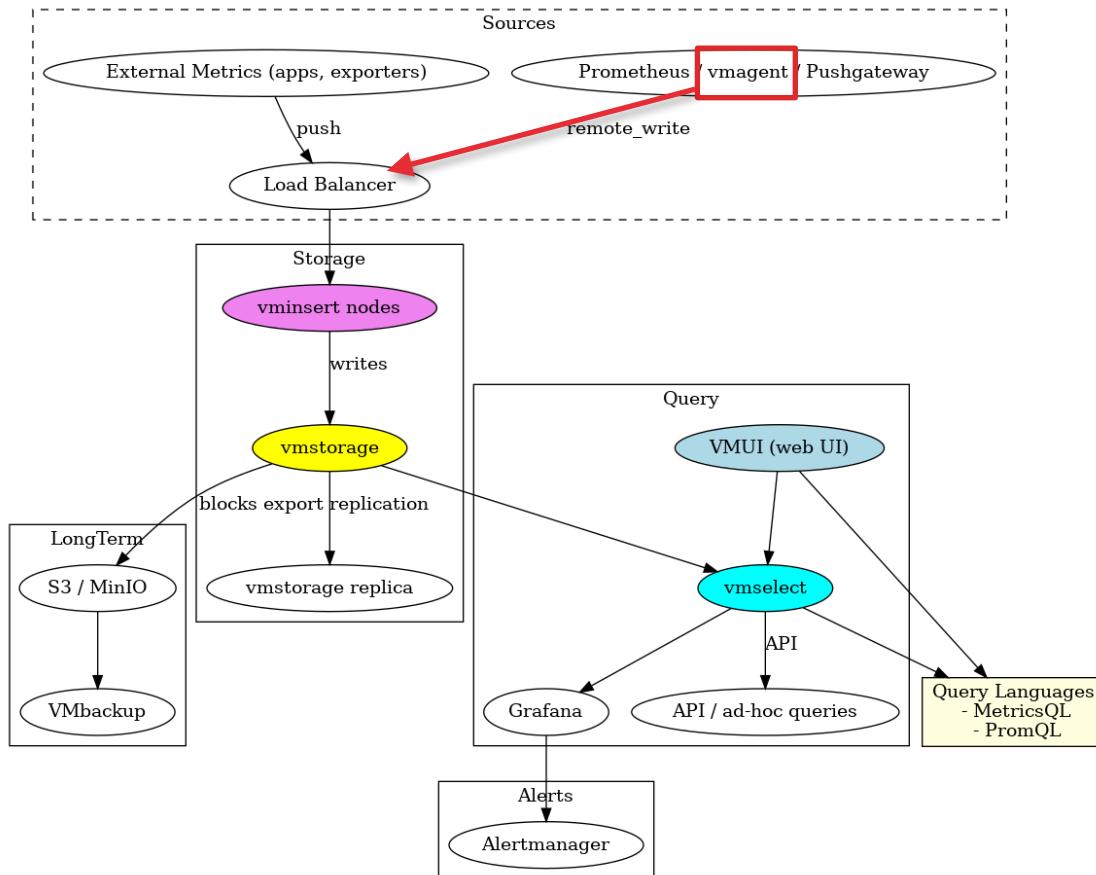
MetricsQL

- ☑ PromQL compatible, no need to change dashboards or alerting rules
- ☑ 100+ additional functions to improve querying experience
- ☑ Templating support in queries

Top 5 features

- ☑ Graphite compatible filters and Graphite support
- ☑ group_left(*) / group_right(*) with prefix support
- ☑ @ modifier for arbitrary timestamps and subexpressions
- ☑ default / if / ifnot binary operators
- ☑ WITH templates + string concatenation





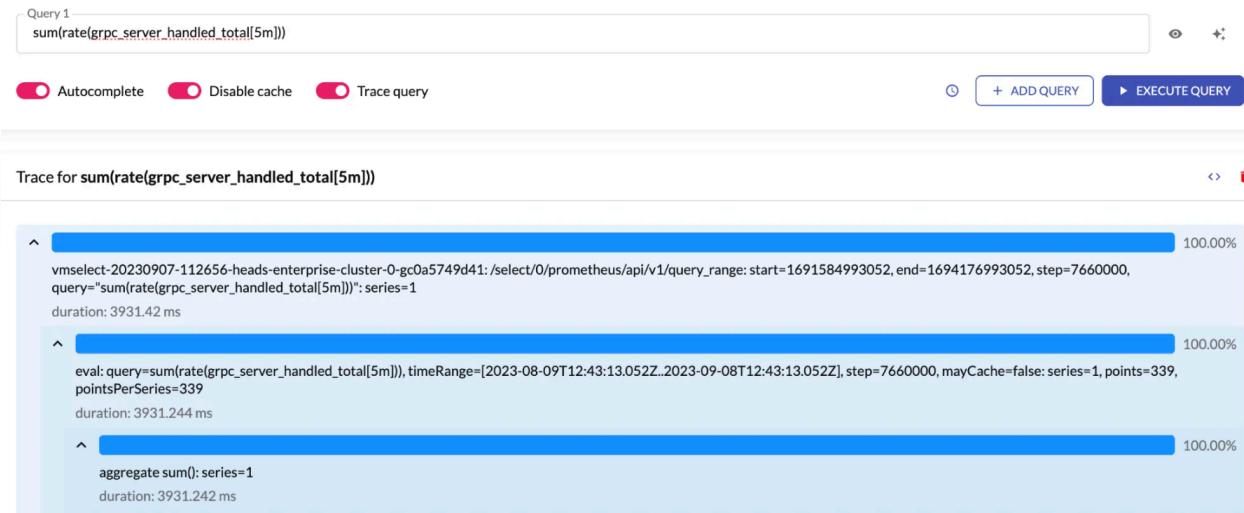
Main gains over Prometheus

Improved
benchmarks

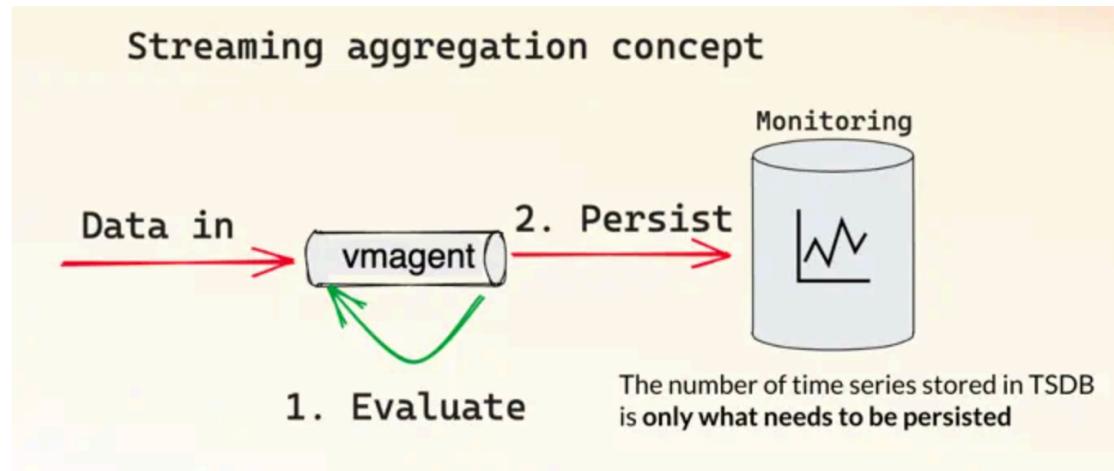
	Prometheus	VictoriaMetrics
CPU avg used	0.79/3 cores	0.76/3 cores
Disk usage	83.5 GiB	✓ 33 GiB
Memory max used	8.12/12 GiB	✓ 4.5/12 GiB
Read latency 50th	70.5ms	✓ 4.3ms
Read latency 99th %ile	7s	✓ 3.6s

Main gains over Prometheus

Query tracing
to find
bottlenecks



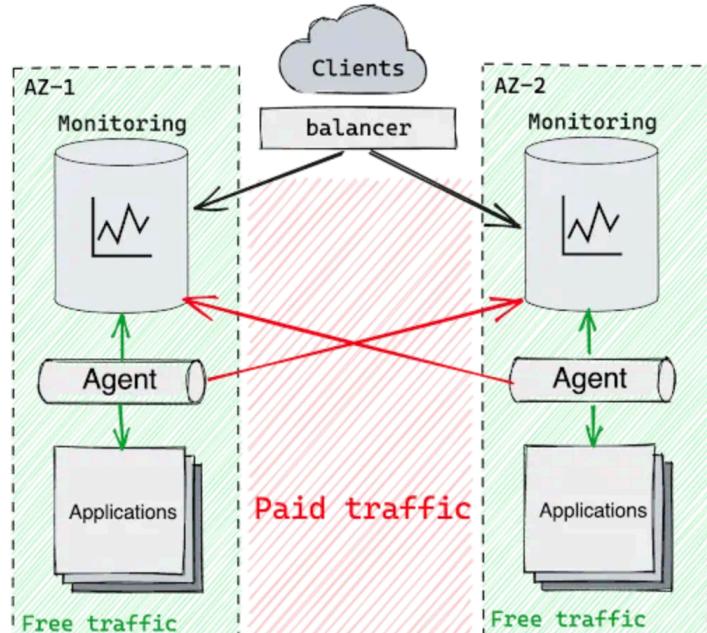
Main gains over Prometheus



Main gains over Prometheus

Save network
cost with
vmagent

Settings	Trade-off
remoteWrite.vmProtoCompressLevel	Increased compression level, traded for higher CPU usage
remoteWrite.maxBlockSize, remoteWrite.maxRowsPerBlock, remoteWrite.flushInterval	Increased batch size leading to better compression ratio, traded for latency
remoteWrite.significantFigures, remoteWrite.roundDigits	Reduced precision/entropy, better compression ratio



Cardinality explorer



Cardinality explorer helps identify

 Metric names with the highest number of series

.....

 Labels with the highest number of series

.....

 Values with the highest number of series for the selected label

.....

 label=name pairs with the highest number of series

.....

 Labels with the highest number of unique values

.....



Spot high cardinality on VictoriaMetrics clusters



High memory usage on vmstorage – indexes for many unique series sit in memory



VictoriaMetrics

Simple, Reliable, Efficient Monitoring

victoriametrics.com

vmstorage



CPU usage spike on vmstorage, each unique series requires a separate TSID mapping and index entry



vmselect - query processing



Memory spikes on vmselect, when queries need to touch millions of series

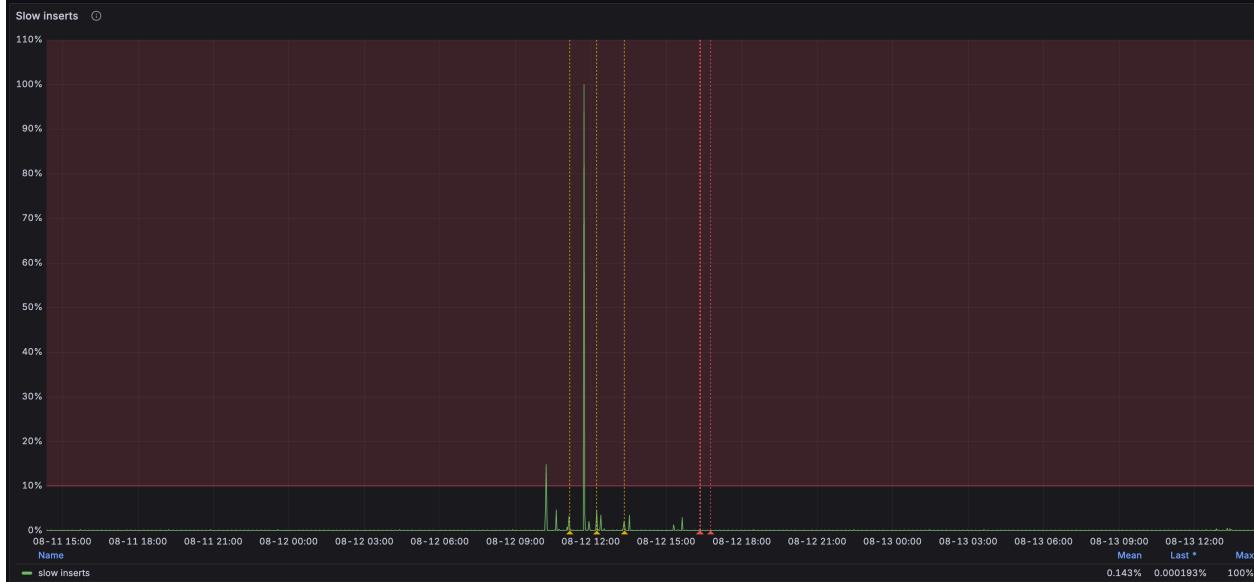


VictoriaMetrics

Simple, Reliable, Efficient Monitoring

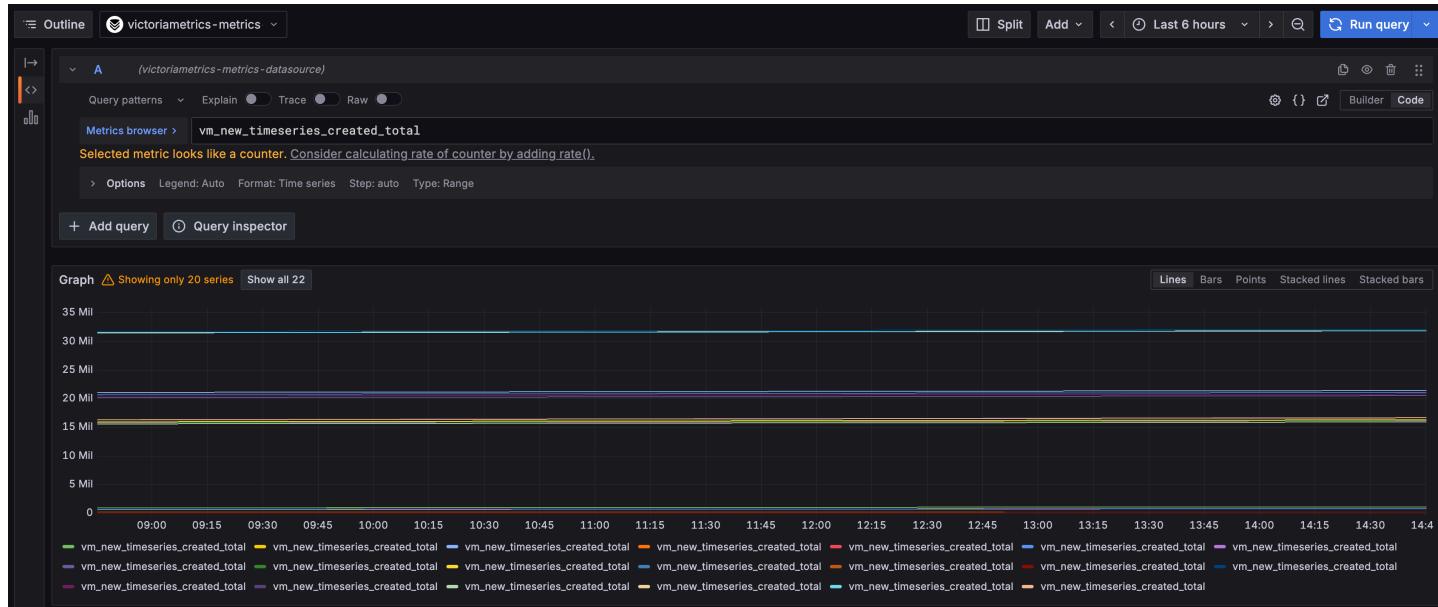
victoriametrics.com

vminsert connects to multiple vmstorage nodes

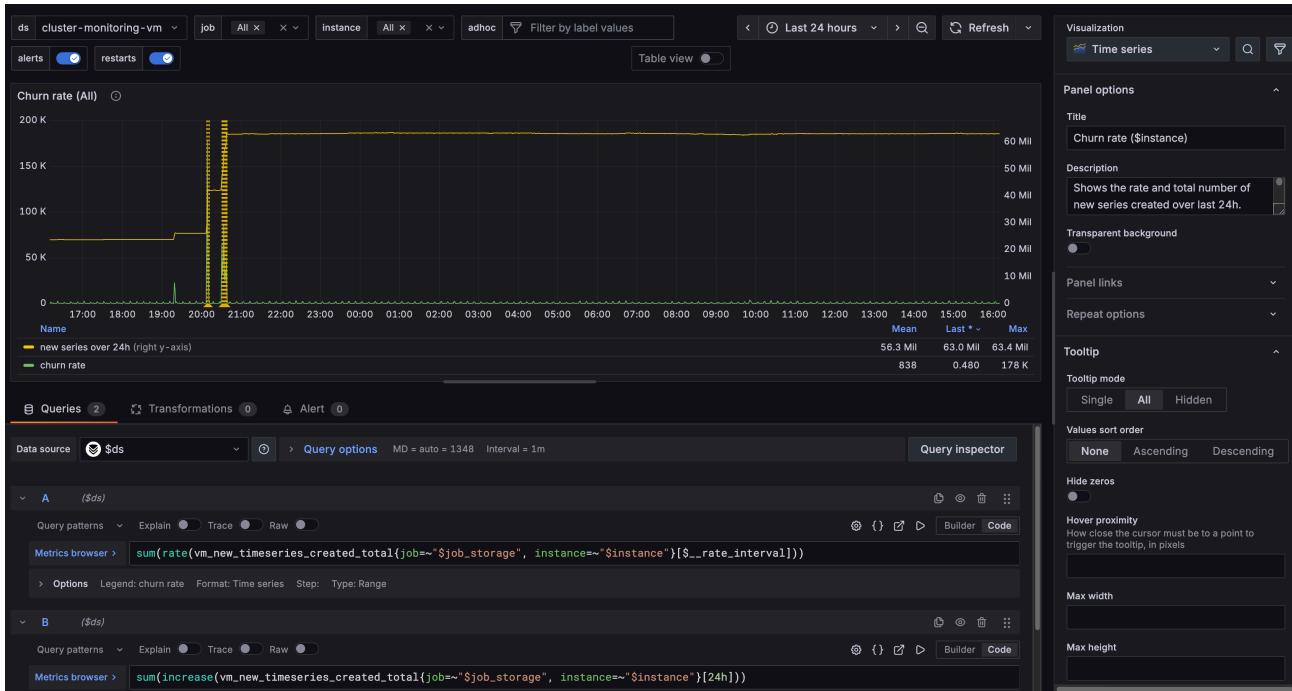


Increased ingestion latency, vminsert takes longer to batch and deduplicate incoming series before writing

Useful queries



How to use it in a Grafana dashboard



Take-aways

- ✓ Reduce histogram buckets, aggregate data with functions like sum() or avg()
- ✓ Look for metrics insights in vmui's Cardinality Explorer
- ✓ If labels can't be removed, pre-aggregate data before ingestion with stream aggregation
- ✓ Reduce query load by aggregating with MetricsQL instead of pulling raw series
- ✓ Correlate your insights with your scrape jobs, dashboards and use cases

Resources



<https://play.victoriametrics.com/> & <https://play-grafana.victoriametrics.com/>



<https://github.com/VictoriaMetrics/prometheus-benchmark>



<https://victoriametrics.com/blog/cardinality-explorer/>



<https://medium.com/@romanhavronenko/victoriametrics-promql-compliance-d4318203f51e>



<https://valyala.medium.com/prometheus-vs-victoriametrics-benchmark-on-node-exporter-metrics-4ca29c75590f>



<https://github.com/VictoriaMetrics-Community/opentelemetry-demo>

More on VictoriaMetrics Solutions



VictoriaMetrics
Open Source Available on Github



VictoriaLogs
Open Source Available on Github



**VictoriaMetrics
Enterprise**

Commercial Licence



**VictoriaMetrics
Cloud**

Pay-as-you-go



Merci d'avoir écouté!

Avez-vous des questions?

Bsky: @didiviking.bsky.social

X: @dianavtodea

Github: @didiViking/Conferences_Talks

LinkedIn: @diana-todea-b2a79968



VictoriaMetrics Community